

Week 5 Assignment - Advanced Supervised Learning

Grad 509: Dimensionality Reduction Assignment

Your task is to develop a Support Vector Machine (SVM) algorithm to solve a real-world problem of identifying the species of Iris Plants. The dataset Module_5_data (see below csv file) is a modified dataset from the original Iris dataset. Your task is to determine whether a given Iris Plant belongs to the Setosa or the Versicolor based on four features: sepal length, sepal width, petal length, and petal width.

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split, GridSearchCV, StratifiedKFold
from sklearn.pipeline import Pipeline
from sklearn.preprocessing import StandardScaler
from sklearn.svm import SVC
from sklearn.metrics import (
    accuracy_score, precision_score, recall_score, f1_score,
    confusion_matrix, ConfusionMatrixDisplay,
    roc_auc_score, RocCurveDisplay, classification_report
)

iris = pd.read_csv('Model_5_data.csv')
```

```
In [2]: iris
```

Out[2]:

| | sepal length (standardized) | sepal width (standardized) | petal length (standardized) | petal width (standardized) | species |
|-----|-----------------------------|----------------------------|-----------------------------|----------------------------|---------|
| 0 | 5.1 | 3.5 | 1.4 | 0.2 | 0 |
| 1 | 4.9 | 3.0 | 1.4 | 0.2 | 0 |
| 2 | 4.7 | 3.2 | 1.3 | 0.2 | 0 |
| 3 | 4.6 | 3.1 | 1.5 | 0.2 | 0 |
| 4 | 5.0 | 3.6 | 1.4 | 0.2 | 0 |
| ... | ... | ... | ... | ... | ... |
| 95 | 5.7 | 3.0 | 4.2 | 1.2 | 1 |
| 96 | 5.7 | 2.9 | 4.2 | 1.3 | 1 |
| 97 | 6.2 | 2.9 | 4.3 | 1.3 | 1 |
| 98 | 5.1 | 2.5 | 3.0 | 1.1 | 1 |
| 99 | 5.7 | 2.8 | 4.1 | 1.3 | 1 |

100 rows × 5 columns

In [7]:

```
X_full = iris.drop('species',axis=1)
y_full = iris['species']

# Filter to a binary problem: Setosa (0) vs Versicolor (1) - just in case there is another species leak
mask = (y_full != 2)
X = X_full[mask]
y = y_full[mask]

print("X shape:", X.shape, "y shape:", y.shape)
print("Class distribution (Setosa=0, Versicolor=1):", {int(c): int((y==c).sum()) for c in np.unique(y)})
```

X shape: (100, 4) y shape: (100,)
 Class distribution (Setosa=0, Versicolor=1): {0: 50, 1: 50}

After doing a review on the internet, an interesting suggestion for feature engineer is to add the ratio between the sepal and petal length with its width.

```
In [8]: # -----
# 2) FEATURE ENGINEERING & TRANSFORMATION
#     - Standardize features (SVMs are scale-sensitive)
#     - Optionally add simple ratios that are sometimes discriminative
# -----
# (Optional) Augment with two simple engineered features:

X = X.to_numpy()
y = y.to_numpy()

sepal_ratio = (X[:, 0] / np.maximum(X[:, 1], 1e-6)).reshape(-1, 1) # sepal_length / sepal_width
petal_ratio = (X[:, 2] / np.maximum(X[:, 3], 1e-6)).reshape(-1, 1) # petal_length / petal_width
X_aug = np.hstack([X, sepal_ratio, petal_ratio])
```

Now before training the model, I will split into to sets (train and test). The split will be stratify so the ratio between Setosa and Versicolor is conserve in both sets

```
In [ ]: # -----
# 3) TRAIN/TEST SPLIT (stratified)
# -----
X_train, X_test, y_train, y_test = train_test_split(
    X_aug, y, test_size=0.30, random_state=24, stratify=y
) # use 70% 30% so we have enough data in test
```

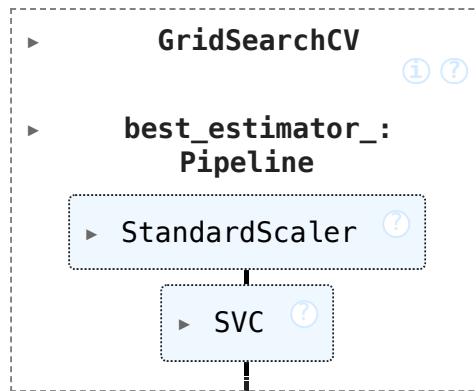
In theory, the data is standarized, based on column name, by as we have add new features, and to be sure I will perform and standarization. Also I will leverage on the **Pipeline class of Scikit-Learn** to handle both standarization and model on the same python object.

```
In [13]: # -----
# 4) MODEL: SVM in a Pipeline
#     - StandardScaler -> SVC
#     - Grid search over linear and RBF kernels
# -----
pipe = Pipeline([
    ("scaler", StandardScaler()),
    ("svc", SVC(probability=True)) # probability=True for ROC; slight overhead
])
```

This part goes beyond what we have seen in the course. But as we can train both SVC using linear separation (classic SVM), or using kernel methods, I will perform a hyperparameter tuning, to see which is the best model.

```
In [14]: # Sensible hyperparameter ranges for this problem:  
# - RBF: keep search compact to avoid overfitting on tiny data.  
param_grid = [  
    {  
        "svc_kernel": ["linear"],  
        "svc_C": [0.01, 0.02, 0.05, 0.1, 0.2]  
    },  
    {  
        "svc_kernel": ["rbf"],  
        "svc_C": [0.01, 0.02, 0.05, 0.1, 0.2],  
        "svc_gamma": ["scale", 0.01, 0.1, 1]  
    }  
]  
  
cv = StratifiedKFold(n_splits=5, shuffle=True, random_state=7)  
grid = GridSearchCV(  
    estimator=pipe,  
    param_grid=param_grid,  
    scoring="roc_auc",  
    cv=cv,  
    n_jobs=-1,  
    refit=True,  
    verbose=0  
)  
grid.fit(X_train, y_train)
```

Out[14]:



```
In [15]: print("\nBest params:", grid.best_params_)
print("CV best score (roc_auc):", f"{grid.best_score_:.4f}")

best_model = grid.best_estimator_
```

```
Best params: {'svc__C': 0.01, 'svc__kernel': 'linear'}
CV best score (roc_auc): 1.0000
```

From model training results, we see that linear is the best choose. And the optimal regularization parameter is 0.01. Now we will see the results on the test set.

In [16]:

```
# -----
# 5) EVALUATION on hold-out test set
# -----

y_pred = best_model.predict(X_test)
y_proba = best_model.predict_proba(X_test)[:, 1]

acc = accuracy_score(y_test, y_pred)
prec = precision_score(y_test, y_pred)
rec = recall_score(y_test, y_pred)
f1 = f1_score(y_test, y_pred)
auc = roc_auc_score(y_test, y_proba)

print("\n==== Test Metrics ===")
print(f"Accuracy : {acc:.4f}")
print(f"Precision: {prec:.4f}")
print(f"Recall   : {rec:.4f}")
```

```
print(f"F1-score : {f1:.4f}")
print(f"ROC-AUC  : {auc:.4f}\n")
print(classification_report(y_test, y_pred, target_names=["Setosa", "Versicolor"]))

# Confusion Matrix
cm = confusion_matrix(y_test, y_pred, labels=[0, 1])
disp = ConfusionMatrixDisplay(confusion_matrix=cm, display_labels=["Setosa", "Versicolor"])
disp.plot(values_format="d")
plt.title("Confusion Matrix – SVM (Setosa vs Versicolor)")
plt.tight_layout()
plt.show()

# ROC Curve
RocCurveDisplay.from_predictions(y_test, y_proba, name="SVM")
plt.title("ROC – SVM (Setosa vs Versicolor)")
plt.tight_layout()
plt.show()
```

==== Test Metrics ===

Accuracy : 1.0000

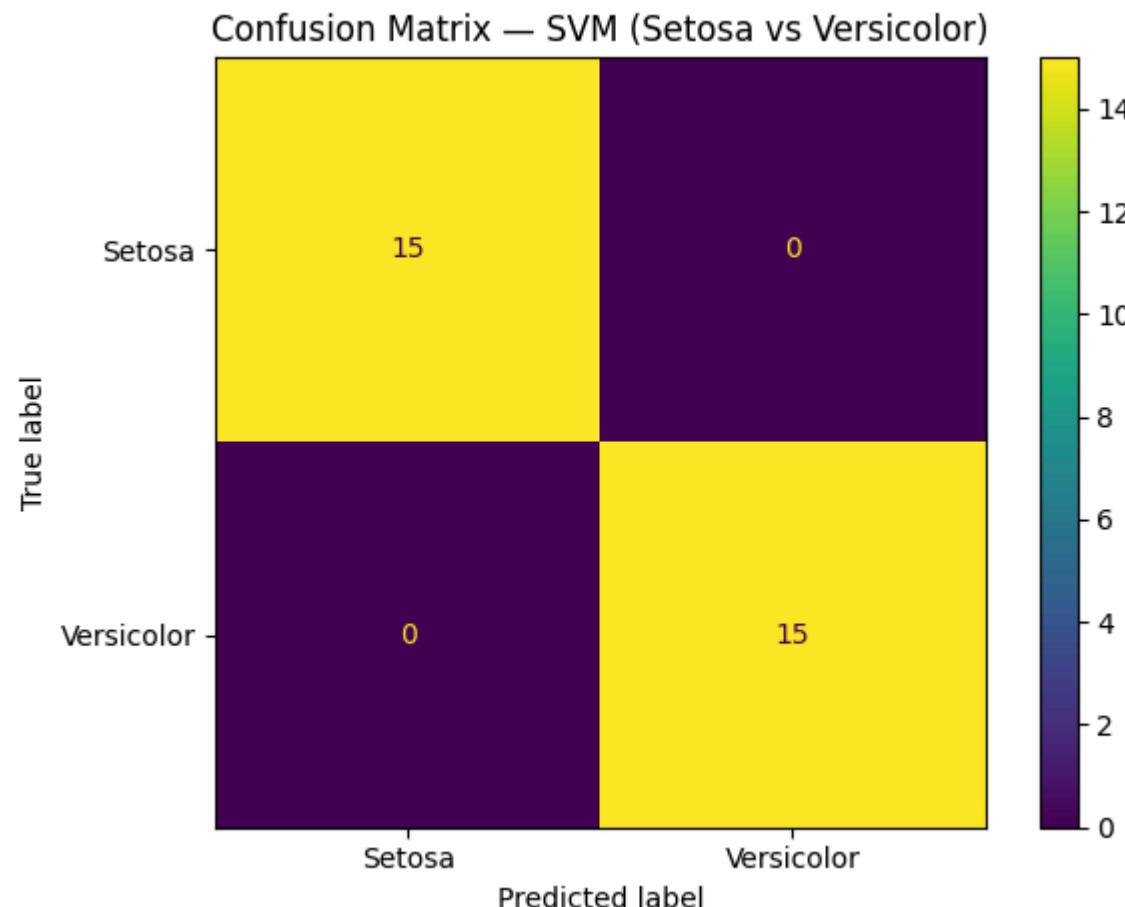
Precision: 1.0000

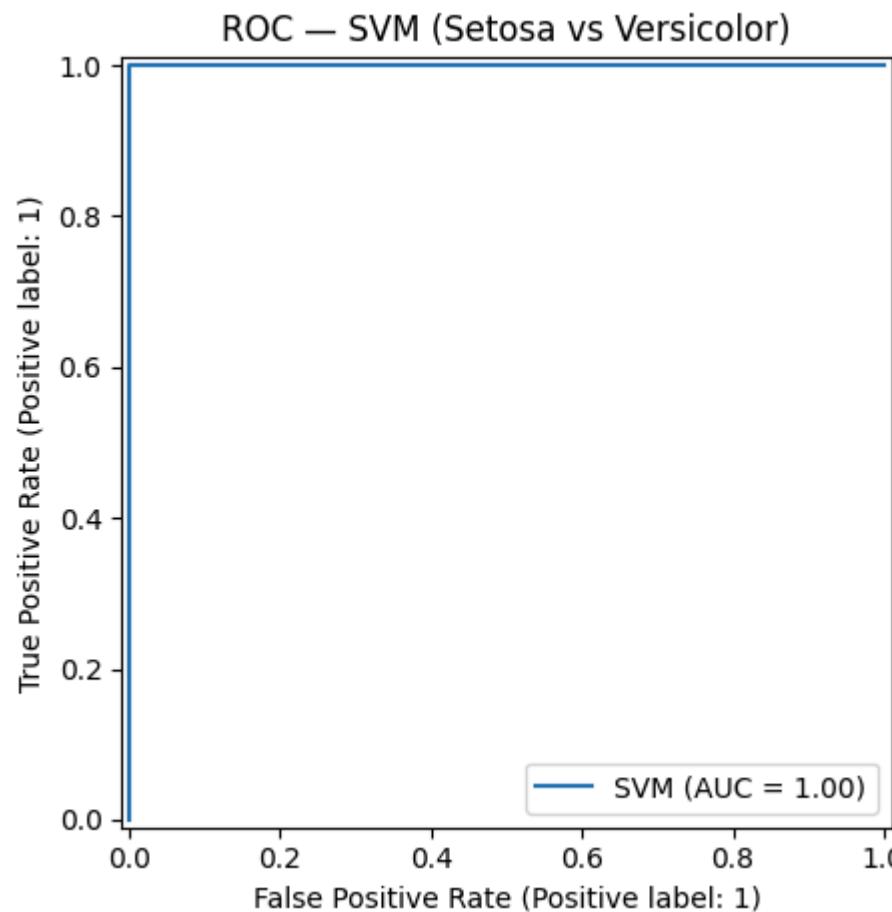
Recall : 1.0000

F1-score : 1.0000

ROC-AUC : 1.0000

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| Setosa | 1.00 | 1.00 | 1.00 | 15 |
| Versicolor | 1.00 | 1.00 | 1.00 | 15 |
| accuracy | | | 1.00 | 30 |
| macro avg | 1.00 | 1.00 | 1.00 | 30 |
| weighted avg | 1.00 | 1.00 | 1.00 | 30 |





We see that the results are impressive. This probably happen because the data is relative simple, and perfectly linear separable:

```
In [20]: feature_name = ['sepal length (standardized)', 'sepal width (standardized)', 'petal length (standardized)', 'petal width (standardized)']
```

```
In [21]: # -----
# 6) INTERPRETABILITY HINT
# -----
if grid.best_params_.get("svc__kernel") == "linear":
    # Access the trained linear SVM inside the pipeline:
    svc = best_model.named_steps["svc"]
    scaler = best_model.named_steps["scaler"]
```

```
# Coefficients are on the scaled feature space. To get a sense of relative importance:  
coefs = svc.coef_.ravel()  
# Pair with feature names:  
for name, w in sorted(zip(feature_name, coefs), key=lambda t: abs(t[1]), reverse=True):  
    print(f"{name:>18s}: {w:+.3f}")
```

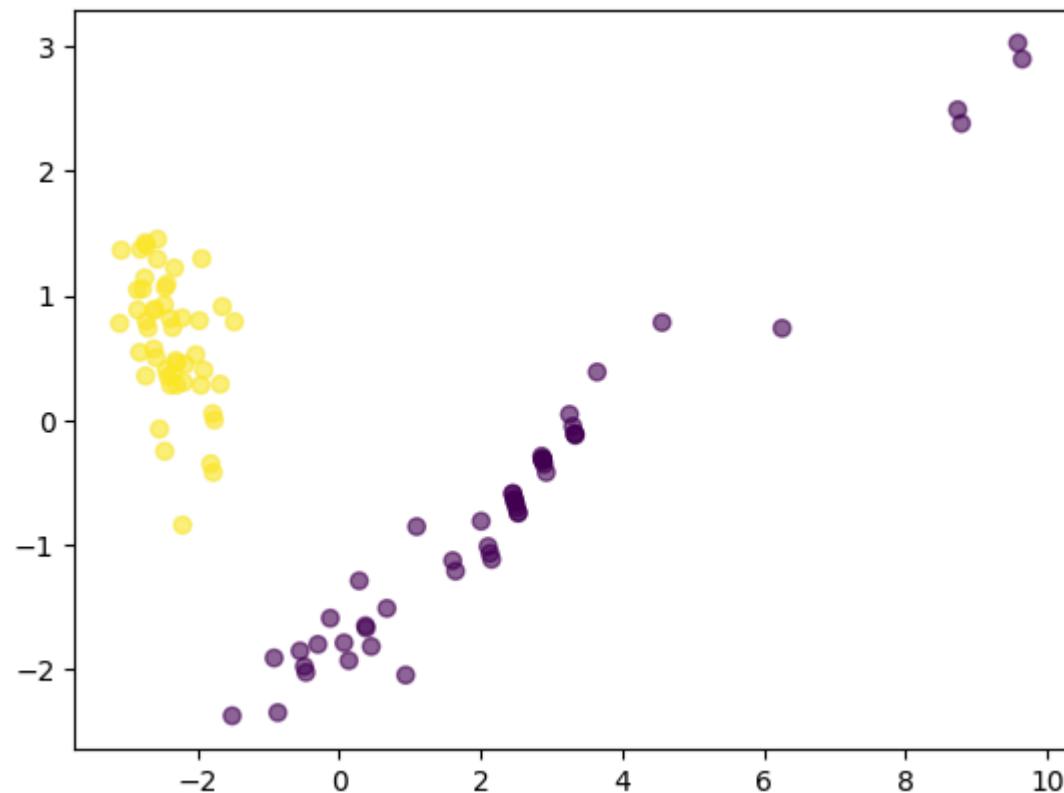
```
petal length (standardized): +0.273  
petal width (standardized): +0.264  
    sepal_ratio: +0.205  
sepal width (standardized): -0.198  
sepal length (standardized): +0.141  
    petal_ratio: -0.122
```

As I have hypothesize, the data must be particularly linearly separable, I will perform a PCA plot, to measure this:

```
In [22]: X_full['sepal_ratio'] = sepal_ratio  
X_full['petal_ratio'] = petal_ratio  
  
X_pca = X_full.to_numpy()  
  
from sklearn.decomposition import PCA  
  
pca = PCA(n_components=2)  
X_reduced = pca.fit_transform(X_pca)
```

```
In [25]: import matplotlib.pyplot as plt  
  
plt.scatter(X_reduced[:,0], X_reduced[:,1], c=y, alpha=0.6)
```

```
Out[25]: <matplotlib.collections.PathCollection at 0x13862e470>
```



From the picture, we confirm that the data is perfectly linearly separable, which explain why our classifier achieve such amazing metrics.